

# Project 1 Readme

## Overview

This README serves as a comprehensive guide to my implementation of a multithreaded file transfer system using the GetFile protocol. The project consists of four progressive components: **Echo** (basic socket communication), **Transfer** (file transfer fundamentals), **GFLib** (GetFile protocol implementation), and **MTGF** (multithreaded server). Each component builds upon the previous, demonstrating mastery of socket programming, protocol design, and concurrent programming with proper synchronization.

## Warmup 1: Echo Client-Server Implementation

### Design Approach and Rationale

**What I Implemented:** A basic client-server pair where the server echoes back any message received from the client (up to 15 bytes).

#### Key Design Decisions:

- **IPv4/IPv6 Dual-Stack Support:** I chose to implement the server using `AF_INET6` with dual-stack capability rather than separate IPv4/IPv6 implementations.
  - **Why:** This approach handles both protocol families with a single socket, reducing complexity
  - **Alternative Considered:** Separate sockets for IPv4 and IPv6, but rejected due to increased resource usage and complexity
- **Client Address Resolution:** Used `AF_UNSPEC` in the client to connect to any available address family.
  - **Why:** Allows the client to connect regardless of server's address family configuration
  - **Implementation:** Iterates through all resolved addresses until successful connection
- Enabling port reuse: `setsockopt(SO_REUSEADDR)`

### Flow of Control - Echo Implementation

ECHO SERVER:

1. `getaddrinfo(NULL, port, AF_INET6, AI_PASSIVE)` → Get server address
2. `socket()` → Create listening socket with dual-stack support
3. `setsockopt(SO_REUSEADDR)` → Enable port reuse for development
4. `bind()` → Bind to address
5. `listen()` → Enter passive listening mode
6. `while (1):`
  - └─ `accept()` → Accept incoming connection
  - └─ `recv(fd, buffer, 16)` → Read client message
  - └─ `send(fd, buffer, received_bytes)` → Echo message back
  - └─ `close(fd)` → Close connection and continue

ECHO CLIENT:

1. `getaddrinfo(hostname, port, AF_UNSPEC)` → Resolve server address
2. `socket()` → Create client socket
3. `connect()` → Attempt connection to server
4. `send(socket, message, length)` → Send message to server
5. `recv(socket, buffer, 16)` → Receive echoed response
6. `fwrite(buffer, received_bytes, stdout)` → Output response
7. `close()` → Clean up and exit

## Testing Strategy - Echo

Tests I Developed:

- **Message Integrity Test:** Verified exact byte-for-byte echo with various message lengths (1-15 bytes)
- **IPv4/IPv6 Compatibility:** Tested server with both IPv4 and IPv6 clients
- **Edge Case Testing:** Empty messages, maximum 15-byte messages, and connection drops

**Debugging Process:** Used CLion with Docker toolchain for development environment consistency

## Warmup 2: Transfer Client-Server Implementation

### Design Approach and Rationale

**What I Implemented:** A file transfer system where the server sends a complete file to any connecting client, and the client saves the received data to disk.

**Critical Design Decision - Partial I/O Handling:**

- **Problem:** Network operations don't guarantee complete data transfer in single calls
- **My Solution:** Implemented comprehensive loops for both send and receive operations

```
// Example from my transfer client - handling partial writes
while (written < received) {
    ssize_t currWritten = write(fd, buffer + written, received - written);
    if (currWritten == -1) handle_error();
    written += currWritten;
}
```

**Why This Approach:** TCP provides a stream abstraction, not message boundaries. The OS may only accept/deliver partial data based on buffer availability.

**Memory Management Strategy:** Used fixed-size buffers (512 bytes) with streaming I/O rather than loading entire files into memory.

- **Why:** Supports large file transfers without memory constraints

- **Alternative Rejected:** Full file buffering - would fail with large files or limited memory

## Flow of Control - Transfer Implementation

TRANSFER SERVER:

1. Initialize socket and bind to port (Same flow as echo)
2. `open(filename, O_RDONLY)` → Open file for reading
3. `while (1):`
  - └─ `accept()` → Accept client connection
  - └─ `lseek(fd, 0, SEEK_SET)` → Reset file pointer to beginning
  - └─ `while (not EOF):`
    - | └─ `read(fd, buffer, BUFSIZE)` → Read file chunk
    - | └─ `while (sent < bytes_read):`
      - | └─ `send(socket, buffer + sent, remaining)` → Send with partial handling
      - | └─ Continue until file complete
    - └─ `close(connection)` → Signal EOF to client

TRANSFER CLIENT:

1. Connect to server
2. `open(output_file, O_WRONLY|O_CREAT|O_TRUNC)` → Create output file
3. `while ((received = recv(socket, buffer, BUFSIZE)) > 0):`
  - └─ `while (written < received):`
    - | └─ `write(fd, buffer + written, remaining)` → Handle partial writes
4. `close(socket)` → Server closes connection to signal completion
5. `close(output_file)` → File transfer complete

## Testing Strategy - Transfer

Tests I Developed:

- **File Integrity Verification:** Used Docker containers to transfer files and verify MD5 checksums match
- **Large File Testing:** Tested with files up to 100MB to verify partial I/O handling
- **Connection Robustness:** Simulated network interruptions during transfer

**Development Environment:** Used CLion with Docker toolchain for development environment consistency

## Part 1: GetFile Protocol Library (GFLib) Implementation

### Design Approach and Rationale

**What I Implemented:** A complete HTTP-like protocol implementation with client library (gfclient) and server library (gfserver) supporting the GetFile protocol.

## Core Design Philosophy - Opaque Pointer Pattern:

```
struct gfcrequest_t {
    char *server;
    unsigned short portno;
    char *path;
    int sfd;
    void (*writefunc)(void *data, size_t data_len, void *arg);
    void *writearg;
    gfstatus_t status;
    size_t fileLength;
    size_t bytesReceived;
};
```

**Why Opaque Pointers:** Provides clean separation between interface and implementation, preventing client code from accessing internal structures while maintaining flexibility for future modifications.

## Protocol State Management Strategy:

- **Challenge:** Parse variable-length headers followed by binary content without excessive buffering
- **My Solution:** Two-phase parsing approach
  1. First phase: Keep receiving header until `\r\n\r\n` delimiter to extract status and content-length
  2. Second phase: Stream body content through registered callbacks

## Callback-Based Architecture:

- **Why:** Allows users to handle data streaming without library making assumptions about storage
- **Benefit:** Memory-efficient for large files - no need to buffer complete file in memory

## Flow of Control - GetFile Protocol

GFCLIENT LIBRARY EXECUTION:

APPLICATION CODE:

```
├─ gfc_create() → Returns opaque gfcrequest_t*
├─ gfc_set_server/port/path() → Configure request parameters
├─ gfc_set_writefunc/writearg() → Register data callback
├─ gfc_perform(request) → Execute request:
|   ├─ socket() and connect() → Establish connection
|   ├─ send("GETFILE GET /path\r\n\r\n") → Send protocol request
|   ├─ Parse response header:
|   |   ├─ recv() until "\r\n\r\n" found → Extract header
|   |   ├─ Parse status (OK/FILE_NOT_FOUND/ERROR/INVALID)
|   |   └─ Extract content-length if status is OK
|   └─ while (bytes_received < content_length):
|       ├─ recv(socket, buffer, chunk_size) → Receive data chunk
|       └─ writefunc(buffer, received, writearg) → Callback to application
```

```
| └─ Return final status to application
└─ gfc_cleanup() → Free all allocated resources
```

#### GFSERVER LIBRARY EXECUTION:

##### MAIN THREAD:

```
└─ gfs_create() → Initialize server structure
└─ gfs_set_handler() → Register request handler callback
└─ gfs_serve() → Start server:
| └─ socket(), bind(), listen() → Setup listening socket
| └─ while (1):
| | └─ accept() → Accept incoming connection
| | └─ Parse request header → Extract method and path
| | └─ Reject invalid and error headers
| | └─ handler(&ctx, path, arg) → Call application handler
| └─ Context cleanup and connection close
```

##### HANDLER CALLBACK (Provided):

```
└─ Validate request path and check file existence
└─ if (file not found): gfs_sendheader(ctx, GF_FILE_NOT_FOUND, 0)
└─ else:
| └─ gfs_sendheader(ctx, GF_OK, file_size) → Send success header
| └─ while (file data remaining):
| | └─ read(fd, buffer, chunk_size) → Read file chunk
| | └─ gfs_send(ctx, buffer, bytes_read) → Send to client
└─ return gfh_success
```

## Testing Strategy - GetFile Protocol

### Tests I Developed:

- **Protocol Compliance Validator:** Custom test scripts sending malformed requests and verifying proper error responses
- **Edge Case Testing:** Invalid schemes, missing headers, incomplete requests
- **Status Code Verification:** Tested all status codes (OK, FILE\_NOT\_FOUND, ERROR, INVALID) with appropriate scenarios

**Development Tools:** CLion with Docker toolchain provided consistent development environment for protocol testing across different network configurations.

## Part 2: Multithreaded GetFile (MTGF) Implementation

### Design Approach and Rationale

**What I Implemented:** A complete multithreaded file server and client using the boss-worker pattern with proper synchronization.

## Critical Design Decision - Boss-Worker Pattern

### Synchronization Strategy:

```
typedef struct {
    int active_workers;
    int shutdown;
    steque_t *queue;
    pthread_mutex_t* mutex;
    pthread_cond_t* worker_cond;
    pthread_cond_t* finish_cond;
} worker_fn_args_t;
```

### Why This Design:

- **Mutex:** Protects shared queue state from race conditions
- **Condition Variables:** Prevent CPU waste when workers are idle - more efficient than busy-waiting
- **Separate Condition Variables:** Different signals for worker wake-up vs. completion notification

### Context Ownership Transfer Challenge:

- **Problem:** Safely passing `gfcontext_t` ownership from main thread to worker threads
- **My Solution:** Ownership transfer pattern with explicit pointer nullification

```
// In handler - transfer ownership to worker
task->ctx = *ctx;
*ctx = NULL; // Prevent double-free by nullifying main thread's pointer
```

## Flow of Control - Multithreaded Server

```
MAIN THREAD (Boss):
├─ pthread_create() × N → Create worker thread pool
├─ steque_init() → Initialize shared work queue
├─ pthread_mutex/cond_init() → Initialize synchronization primitives
├─ while (server running):
│   └─ accept() → Wait for client connection (blocking)
│   └─ Create task_item_t:
│       └─ task->ctx = gfcontext (transfer ownership)
│       └─ task->path = requested_path
│       └─ task->arg = application_argument
│   └─ pthread_mutex_lock(mutex)
│   └─ steque_push(queue, task) → Add work to queue
│   └─ pthread_cond_signal(worker_cond) → Wake up waiting worker
│   └─ pthread_mutex_unlock(mutex)
│   └─ Continue to next accept() (connection now owned by worker)
```

WORKER THREADS (N workers running worker\_fn):

```
┌ while (1): // Infinite worker loop
|   ┌ pthread_mutex_lock(mutex)
|   ┌ while (steque_isempty(queue)):
|     ┌ pthread_cond_wait(worker_cond, mutex) → Sleep until work available
|   ┌ task = steque_pop(queue) → Get work item
|   ┌ pthread_mutex_unlock(mutex)
|   ┌ content_get(task→path) → Open requested file (fd or -1)
|   ┌ if (fd == -1):
|     ┌ gfs_sendheader(task→ctx, GF_FILE_NOT_FOUND, 0)
|   ┌ else:
|     ┌ fstat(fd, &file_stat) → Get file size
|     ┌ gfs_sendheader(task→ctx, GF_OK, file_size)
|     ┌ while (offset < file_size):
|       ┌ pread(fd, buffer, chunk_size, offset) → Read file chunk
|       ┌ gfs_send(task→ctx, buffer, bytes_read) → Send to client
|       ┌ offset += bytes_read
|     ┌ close(fd)
|   ┌ gfs_cleanup(task→ctx) → Clean up connection resources
|   ┌ free(task) → Free work item and continue loop
```

CRITICAL SYNCHRONIZATION POINTS:

- Queue access: All steque operations protected by mutex
- Worker coordination: Condition variables prevent busy-waiting
- Context handoff: Ownership transfer prevents double-free errors
- Resource cleanup: Each worker responsible for complete cleanup

## Multithreaded Client Implementation

CLIENT MAIN THREAD (Boss):

```
┌ Parse workload file → Generate list of download requests
┌ steque_push(queue, request) → Push all download requests to queue
┌ pthread_create() × N → Create worker thread pool
┌ pthread_mutex_lock(mutex)
┌ pthread_cond_broadcast(worker_cond) → Wake all workers
┌ while (steque_isempty(queue) || active_workers > 0):
|   ┌ pthread_cond_wait(finish_cond, mutex) → Wait for completion
┌ pthread_join() → Clean up worker threads
┌ clean up mutex/cond/queue/tids
```

#### CLIENT WORKER THREADS:

```
|— while (1):
|   |— pthread_mutex_lock(mutex)
|   |— while (steque_isempty(queue) && !shutdown):
|       |— pthread_cond_wait(worker_cond, mutex)
|   |— if (steque_isempty(queue) && shutdown):
|       |— break → Exit worker loop
|   |— request = steque_pop(queue)
|   |— active_workers++ → Increment active count
|   |— pthread_mutex_unlock(mutex)
|   |— gfc_perform(request) → Execute download using gfcclient library
|   |— pthread_mutex_lock(mutex)
|   |— active_workers-- → Decrement active count
|   |— if (steque_isempty(queue) && active_workers == 0):
|       |— pthread_cond_signal(finish_cond) → Signal potential completion
|   |— pthread_mutex_unlock(mutex)
```

## Testing Strategy - Multithreaded Implementation

### Tests I Developed:

- **Concurrency Stress Testing:** Created custom workloads with 100+ simultaneous requests to identify race conditions
- **Thread Pool Scaling:** Tested with 1, 4, 8, 16, 32 worker threads to find optimal configuration
- **Resource Leak Detection:** Extended operation tests to verify no memory leaks or file descriptor leaks
- **Deadlock Prevention:** Systematic testing of different request orderings and timing scenarios

## References - Code and Materials Used

### Code I Used But Did Not Write

- **Steque Library (steque.c/steque.h):** Used the provided thread-safe queue implementation for work distribution between boss and worker threads
- **Content Library (content.c/content.h):** Used provided file access abstraction for server-side file operations
- **Workload Library (workload.c/workload.h):** Used provided workload parsing for client request generation

### Code I Copied Within My Own Project

- **Socket Setup Pattern:** Replicated the `getaddrinfo()/socket()/bind()` sequence from echo server in transfer server and gfserver implementations



- **Error Handling Patterns:** Consistent error checking and cleanup patterns copied across all components
- **Partial I/O Loops:** Similar send/receive loop patterns used in transfer, gfclient, and gfserver implementations

## Technical References I Consulted

- **POSIX Threads Programming Tutorial (LLNL):** For understanding condition variable usage patterns and avoiding common deadlock scenarios
- **Beej's Guide to Network Programming:** Referenced for socket programming best practices and IPv4/IPv6 dual-stack implementation
- **CS6200 Lecture Slides on Multithreading:** Boss-worker pattern design and synchronization primitive selection
- **Stack Overflow Thread #23207312:** "pthread condition variable example" - helped understand proper condition variable signaling patterns
- **GNU C Library Manual:** File I/O operations documentation, specifically for pread() usage and file descriptor management
- **Course Piazza Posts:** Specific clarifications about GetFile protocol format and threading requirements

## Development Tools I Used

- **CLion IDE with Docker Toolchain:** Primary development environment with integrated debugging support
- **Docker:** Containerized development environment for consistent testing across platforms

**Note:** All external materials were used for concept understanding and best practice guidance. No code was directly copied from external sources - implementations represent my own understanding and application of these concepts.

---

## Project Improvement Suggestions

### 1. Enhanced Local Testing Framework

**Problem I Observed:** Students lack comprehensive testing tools and rely heavily on Gradescope for validation, limiting learning opportunities.

**Actionable Solution:** Create a local testing framework that provides immediate feedback

### 2. Improved Project Instructions

**Problem I Encountered:** The GFLib implementation instructions are unclear about responsibility boundaries between student code and provided callbacks.

**Current Instructions Say:** "The handler callback should not contain any of the socket-level code. Instead, it should use the gfs\_sendheader, gfs\_send, and potentially the gfs\_abort functions

provided by the gfserver library."

**What's Confusing:** I don't know whether invalid and wrong headers should be handled by my logic or the provided handler callback.

Actionable Solution: Include a decision flowchart showing exactly when student protocol code vs. application handler should handle different error conditions.